

Chatbot Lab

Overview

In order to determine how various prompts (instructions) impact the interaction, three different versions of a chatbot that takes pizza orders were evaluated. The three types of prompts were intended to direct the chatbot's behavior: Normal, Friendly, and Efficient. Every version has distinct objectives:

Normal: An order-taking bot that is clear and succinct.

Friendly: A humorous, amiable, and engaging helper.

Efficient: A quick and simple bot that reduces pointless interaction.

Methodology

To test the variations, the chatbot was implemented with each of the three prompt styles, and different pizza-related inquiries were made by the user. Each version aimed to fulfill the order while adhering to its specific behavior:

Normal: The bot collected all the information required for the order and verified everything before ending the chat in a clear and professional manner.

Friendly: The bot kept the conversation light, offered food suggestions, and engaged the user with humor in a more relaxed manner.

Efficient: By asking only the most important questions, the bot concentrated on rapidly obtaining the order data with the least amount of back and forth.

Issues Observed:

Hallucinations/Inaccurate Reactions The bot occasionally produced unnecessary or superfluous information in the Friendly version.

When asked about drink preferences, for instance, it might provide recommendations unrelated to the user's selections, which could cause some confusion. This occurred as a result of the bot becoming overly informal and deviating from receiving orders.

Unintentional Outcomes: In the Efficient version, the bot occasionally omitted important details, including verifying drink volumes or topping selections. Although the bot tried to be brief, it left out details that might have caused mistakes in the final order.